

Chapter 21

stevenjin8

June 28, 2021

Exercises

Exercise 1

This exercise confused me very much because I thought it was asking to use Laplace distribution to approximate the posterior, rather than a Laplacian approximation of the posterior. Following the presentation in section 8.4.1, the posterior is given by

$$p(\mu, \ell | \mathcal{D}) = \frac{1}{Z} e^{-E(\mu, \ell)} \quad (1)$$

where

$$E(\mu, \ell) = n \log \sigma + \frac{n}{2\sigma^2} (s^2 + (\bar{x} - \mu)^2)$$

which is just equation 21.203 negated.

Finding the gradient is straightforward:

$$\begin{aligned} \frac{\partial E}{\partial \mu} &= -\frac{n(\bar{x} - \mu)}{\sigma^2} \\ \frac{\partial E}{\partial \ell} &= \frac{\partial}{\partial \ell} \left[n\ell + e^{-2\ell} \frac{n}{2} (s^2 + (\bar{x} - \mu)^2) \right] \\ &= n - e^{-2\ell} n (s^2 + (\bar{x} - \mu)^2) \\ &= n - \frac{n (s^2 + (\bar{x} - \mu)^2)}{\sigma^2}. \end{aligned}$$




Figure 1: True log posterior (red) vs Laplacian approximation (blue). Data are 15 samples from a unit normal distribution ($\mu, \ell = 0$).

The Hessian is also pretty easy to compute:

$$\begin{aligned}\frac{\partial^2 E}{\partial \mu^2} &= \frac{n\bar{x}}{\sigma^2} \\ \frac{\partial^2 E}{\partial \mu \partial \ell} &= \frac{2n(\bar{x} - \mu)}{\sigma^2} \\ \frac{\partial^2 E}{\partial \ell^2} &= \frac{\partial}{\partial \ell} n - e^{-2\ell} n (s^2 + (\bar{x} - \mu)^2) \\ &= 2 \frac{n (s^2 + (\bar{x} - \mu)^2)}{\sigma^2}.\end{aligned}$$

Setting the first derivatives (i.e. the gradient) to 0 and solving gives

$$\mu^* = \bar{x}, \ell^* = \log s$$

Thus, the Hessian of E at $\mu = \mu^*, \ell = \ell^*$ is

$$\mathbf{H} = \begin{bmatrix} \frac{n}{s^2} & 0 \\ 0 & 2n \end{bmatrix}$$

We can approximate E with

$$E(\boldsymbol{\theta}) \approx E(\boldsymbol{\theta}^*) + (\boldsymbol{\theta} - \boldsymbol{\theta}^*)^T \mathbf{H}(\boldsymbol{\theta} - \boldsymbol{\theta}^*)$$

where $\boldsymbol{\theta} = (\mu, \ell)$. Finally, we plug this into equation 1:

$$\begin{aligned}p(\boldsymbol{\theta}|\mathcal{D}) &\propto e^{E(\boldsymbol{\theta}^*) + (\boldsymbol{\theta} - \boldsymbol{\theta}^*)^T \mathbf{H}(\boldsymbol{\theta} - \boldsymbol{\theta}^*)} \\ &\propto e^{(\boldsymbol{\theta} - \boldsymbol{\theta}^*)^T \mathbf{H}(\boldsymbol{\theta} - \boldsymbol{\theta}^*)}\end{aligned}$$

We can see that

$$p(\mu, \ell) \approx \mathcal{N}(\boldsymbol{\theta} | \boldsymbol{\theta}^*, \mathbf{H}^{-1}).$$

See Figure 1 for a visualization.

I found the way the author suggested to this problem to be a bit unintuitive because it makes you think that the Hessian in equation 21.206 is the inverse of the covariance matrix. However, this cannot be true since the entries on the diagonal would be negative.

Exercise 6

The goal here is to minimize

$$\mathbb{KL}(q_i || \tilde{p}_i) = \sum_{x \in \{-1, 1\}} q_i(x) \log \frac{q_i(x)}{\tilde{p}_i(x)}$$

This can be achieved by setting $q_i(x) \propto \tilde{p}_i(x)$. Thus, we have

$$q_i(x = 1) = \frac{e^{m_i + L^+}}{e^{m_i + L^+} + e^{-m_i + L^-}}$$

where m_i, L^+ , and L^- have the same meanings as in section 21.3.2. The result follows from equations 21.46-21.50.

Exercise 7

$$\begin{aligned} \mathbb{KL}(p || q) &= \sum_{x, y} p(x, y) \log \frac{p(x, y)}{q(x)q(y)} \\ &= \sum_{x, y} p(x, y) (\log p(x, y) - \log q(x) - \log q(y)) \\ &= \sum_{x, y} p(x, y) \log p(x, y) - \sum_y p(y) \log q(y) - \sum_x p(x) \log q(x) \end{aligned}$$

We minimize this by setting $q(x) = p(x)$ and $q(y) = p(y)$.

Following sections 21.2.2, $p(x, y) = 0 \implies q(x)q(y) = 0$. It follows that the three local minima are

$$\begin{aligned} q(x) = q(y) &= [0.5, 0.5, 0, 0] \\ q(x) = q(y) &= [0, 0, 1, 0] \\ q(x) = q(y) &= [0, 0, 0, 1] \end{aligned}$$

Which have a KL divergence of $\log 2$, $\log 4$, and $\log 4$ respectively.

By setting $q(x, y) = p(x)p(y)$, the KL divergence diverges.